

Stable Hot Spot Analysis (Draft)

Marc Gassenschmidt, Viliam Simko, and Julian Bruns
todo@fzi.de, simko@fzi.de, bruns@fzi.de

FZI Forschungszentrum Informatik
am Karlsruher Institut für Technologie
76131, Haid-und-Neu-Str. 10-14
Karlsruhe, Germany

Abstract: Hot spot analysis is essential for geo-statistics. It supports decision making by detecting points as well as areas of interest in comparison to their neighbourhood. However, these methods are dependent on different parameters, ranging from the resolution of the study area to the size of their neighbourhood. This dependence can lead to instabilities of the detected hotspots, where the results can highly vary between different parameters. A decision maker can therefore ask how valid the analysis actually is. In this study, we examine the impact of key parameters on the stability of the hotspots, namely the size of the neighbourhood, the resolution and the size of the study area, as well as the influence of the ratio between those parameters. We compute the hotspots with the well known Getis-Ord (G^*) statistic as well as its modification, the *Focal G^** statistic. We measure the stability of the hotspot analysis using a recently introduced *stability of hotspots* metric (SoH) and compare the results to intuitive visual analysis. We evaluate the results on real world data with the well-known yellow cab taxi data set from New York, Manhattan. Our results indicate a negative impact on the stability with an increase of the size of the neighbourhood as well as a reduction of the size of the study area, regardless of the resolution.

1 Introduction

The goal of hotspot analysis is the detection and identification of interesting areas. It achieves this goal by computing statistically significant deviations from the mean value of a given study area. This allows a decision maker to easily identify those areas of interest and allows further focus in sub-sequential data analysis or the decision focus. Typical applications range from crime detection over identification of disease outbreaks to urban heat islands. In such applications, scarce resources are then often applied in only those identified hotspots or used as the basis for the allocation. The general approach is an unsupervised learning method similar to a cluster analysis.

But, similar to a cluster analysis, there does exist a high dependency of the identified hotspots on the detection method and in particular the parametrization of this method. The identified areas as well as their shape can vary highly. This volatility can lead to a decrease in trust in the result or in suboptimal allocations of scarce resources. Therefore it is necessary to measure and evaluate

the stability of a hotspot analysis as well as the different parametrizations. In our initial work (Hier Bruns and Viliam 2017), we introduced a method to measure the stability of hotspots, the *stability of hotspots* metric (SoH) and showed its use on the basis of temperature data. Here, we build upon that work and examine in more detail the impact of the different instantiations of the most typical parameter. We use the well known Getis-Ord statistic [?], the standard G^* , and a modification of this statistic, the focal G^* (cite bruns and viliam). Those parameter are the size of the study area (focal matrix), the detail of the resolution (zoom) and the size of the neighbourhood (weight matrix). By varying over these parameter, we can compare the stability for all possible combinations and isolate the effect of single parameter by aggregation over the other parameter. We evaluate on the well-known taxi data set to show the applicability on real world use cases as well as to enable a simple replication.

2 Related Work

2.1 Quality of Clustering

The problem of assessing the quality in unsupervised learning is well known. In the case of the k-mean algorithm, the quality of the clustering is mostly dependent on the value of the k and a miss-specification can lead to highly irregular clusters. In a simple 2D clustering, they can be easily recognized by visual analysis, but in higher dimensionality, this is impossible. One method, to measure the quality of such a clustering is the compactness of the clusters, see e.g. Song 2010 (<http://ieeexplore.ieee.org/document/5558958/?part=1>). This enables the comparison between different clusters. Another possibility is the Silhouette Coefficient by Kaufman et Rousseeuw 1990. This metric measures the similarity of objects in a cluster in comparison to other clusters. For density based clustering, e.g. for DBSCAN, OPTICS (Ankerst et al 1999) gives a simple method to tune the essential parameter for this clustering. This is only a small overview of methods to influence and measure the quality of different clustering methods. But it shows that this problem is not easily solved and dependent on the chosen algorithm. To our knowledge, there does not exist a method to overall measure the stability of a clustering.

2.2 Hot Spot Analysis

The goal of hotspot analysis is the detection of interesting areas as well as patterns in spatial information. One of the most fundamental approach is Moran's I [?]. There it is tested whether or not a spatial dependency exists. This gives the information on global dependencies in a data set. Upon this hypothesis test several geo-statistical tests are based. The most well known are the Getis-Ord statistic [?] and LISA [?]. In both cases the general, the global statistic of Moran's I is applied in a local context. The goal is to detect not only global values, but instead to focus on local hotspots and to measure the significance of those local areas. A more in depth overview of methods to identify and visualize spatial patterns and areas of interest can be found in Shekhar et al (2011).

3 Stability of Hot Spot Analysis

Existing methods for determining hot spots are dependent on the parametrization of the weight matrix as well as on the size of the study area. Intuitively, increasing the size of a weight matrix has a "blurring" effect on the raster (Fig. ??a) whereas decreasing the size can be seen as a form of "sharpening" (Fig. ??b).

For a data analyst, when exploring the data interactively by choosing different filter sizes (weight matrices) or point aggregation strategies (pixel sizes), it is important that the position and size of a hot spot changes in a predictable manner. We formalize the intuition in our stability metric.

We define a hot spot found in comparably more coarse resolutions as *parent* (larger weight matrix or larger pixel size) and in finer resolutions as *child* (smaller weight matrix or smaller pixel size).

That is

To be stable, one assumes that every parent has at least one child and that each child has one parent. For a perfectly stable interaction, it can be easily seen that the connection between parent and child is a injective function and between child and parent a surjective function. To measure the closeness of connection, we propose a metric called the *Stability of Hot spot* (SoH). It measures the deviation from a perfectly stable transformation of resolutions.

In its downward property (from parent to child, injective) it is defined as:

$$SoH^\downarrow = \frac{ParentsWithChildNodes}{Parents} = \frac{|Parents \cap Children|}{|Parents|} \quad (1)$$

And for its upward property (from child to parent, surjective):

$$SoH^\uparrow = \frac{ChildrenWithParent}{Children} = 1 - \frac{|Children - Parents|}{|Children|} \quad (2)$$

where *ParentsWithChildNodes* is the number of parents that have at least one *child*, *Parents* is the total number of

parent, *ChildrenWithParent* is the number of children and *Children* as the total number of children. The SoH is defined for a range between 0 and 1, where 1 represents a perfectly stable transformation while 0 would be a transformation with no stability at all.

4 Focal Getis-Ord

4.1 Dataset

The two datasets (morning and evening flights) depicted in Fig. ?? and Fig. ?? were obtained from a thermal flight over the city of Karlsruhe on 26.09.2008 at 6:30–7:45 and 20:00–21:30. The flights were executed by the Nachbarschaftsverband Karlsruhe¹. A single pixel in the raster represents an area of approximately $5 \times 5m$. The whole dataset of size $35 \times 25km$ was cropped into the inner city area of $2.4 \times 1.4km$. The temperatures in our dataset range from $-1.7^\circ C$ to $18.3^\circ C$. Missing values in the dataset were interpolated using a focal median function with a square matrix of 11×11 pixels, mainly for speeding up further computations and to avoid special handling of NA values.

4.2 Method

In the following text, we use the notation $R \overset{op}{\circ} M$ to denote a focal operation *op* applied on a raster *R* with a focal window determined by a matrix *M*. This is roughly equivalent to a command `focal(x=R, w=M, fun=op)` from package *raster* in the R programming language [?].

Definition 1 (G^* function on rasters). *The function G^* can be expressed as a raster operation:*

$$G^*(R, W, st) = \frac{R \overset{sum}{\circ} W - M * \sum_{w \in W} w}{S \sqrt{\frac{N * \sum_{w \in W} w^2 - (\sum_{w \in W} w)^2}{N-1}}}$$

where:

- *R* is the input raster.
- *W* is a weight matrix of values between 0 and 1.
- *st* = (*N*, *M*, *S*) is a parametrization specific to a particular version of the G^* function. (Def. 2 and 3).

Definition 2 (Standard G^* parametrization). *Computes the parametrization *st* as global statistics for all pixels in the raster *R*:*

- *N* represents the number of all pixels in *R*.
- *M* represents the global mean of *R*.
- *S* represents the global standard deviation of all pixels in *R*.

¹<http://www.nachbarschaftsverband-karlsruhe.de/>

Definition 3 (Focal G^* parametrization). *Let F be a boolean matrix such that: $\text{all}(\text{dim}(F) \geq \text{dim}(W))$. This version uses focal operations to compute per-pixel statistics given by the focal neighbourhood F as follows:*

- N is a raster computed as a focal operation $R \overset{\text{sum}}{\circ} F$. Each pixel represents the number of pixels from R convoluted with the matrix F .
- M is a raster computed as a focal mean $R \overset{\text{mean}}{\circ} F$, thus each pixel represents a mean value of its F -neighbourhood.
- S is a raster computed as a focal standard deviation $R \overset{\text{sd}}{\circ} F$, thus each pixel represents a standard deviation of its F -neighbourhood.

5 Results

5.1 Heatmap

6 Evaluation

- To evaluate, we compare G^* with Focal G^* on the same dataset. - We use NY taxi dropoffs - a single evaluation run is defined in Def. 4. - In an ideal case, we could produce a 3D-plot: - x-axis would be growing pixel sizes 100..1000 by 100 (aka zoom-out) - y-axis would be growing weight matrix sizes (e.g. 5..41 by 2) - z-axis would be performance of G^* vs Focal G^* x and y coordinates - we could also repeat the computation for multiple datasets to obtain multiple samples and to compute error bars. - due to paper size limit, we choose projection in which matrix size is set to constant 7x7 pixels and the only variable is the pixel size. This way, we produce a 2D plot. We also generate just a single sample.

Definition 4 (Evaluation Run). *We define a single evaluation run as a tuple:*

$$E = (V, m, p, w)$$

where:

- V is the input dataset of points, representing the taxi dropoffs in our case.
- m is the metric used, in our case either SoH^\uparrow or SoH^\downarrow .
- p represents the pixel size for aggregating points from V , e.g. 100×100 meters.
- w represents the size of a weight matrix. In our case, we chose a weight matrix depicted in Figure ??(a) for both the G^* and Focal G^* cases.

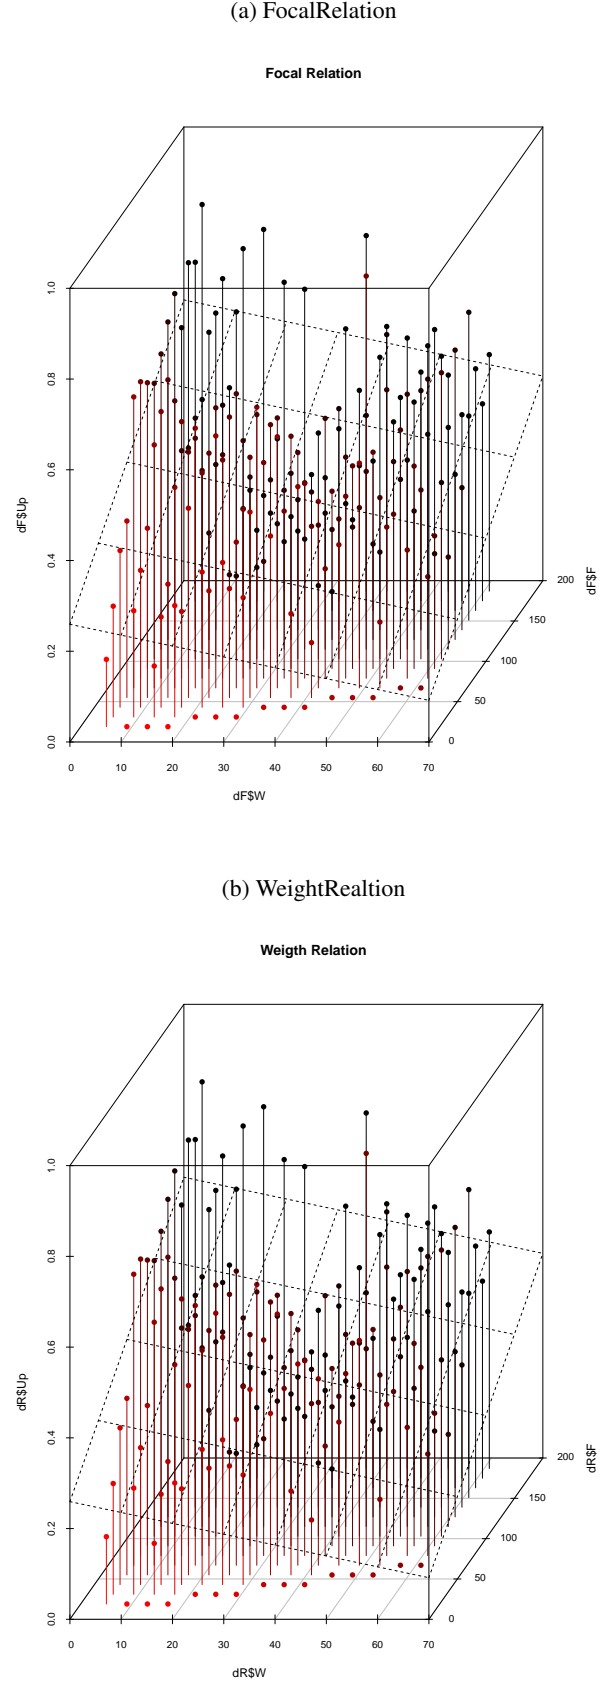




Figure 1: Heatmap of Relation between size of weight matrix W and focal matrix F

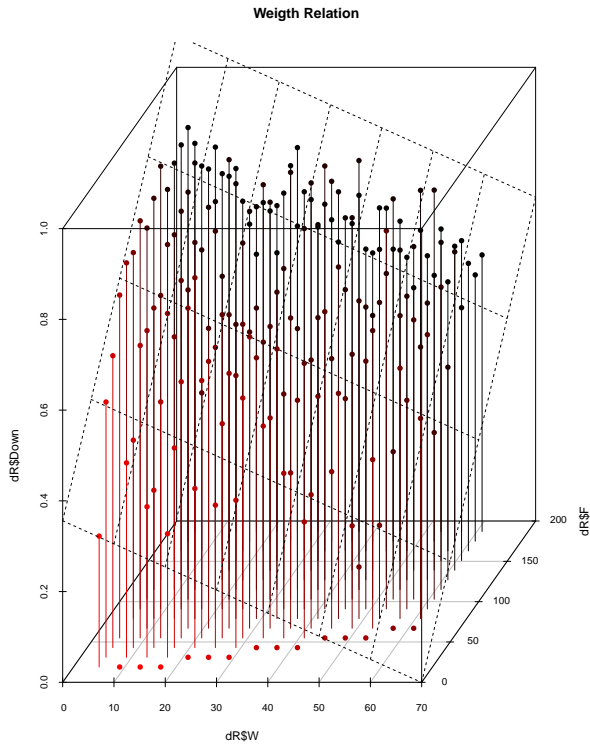
(a) FocalRelation



(a) Upward



(b) WeightRelation



(b) Downward




Figure 2: Heatmap of Relation between size of weight matrix W and focal matrix F

Figure 3: TODO

6.1 Clumping

6.2 Zoom

Blueline is Focal G^* . Redline is G^*

6.3 Blur

The evaluation results are plotted in Fig. ??, each point in the graph represents the SoH^\uparrow metric (Eq.2) between two G^* generated using weight matrices of size i and $i+2$. The focal matrix F has a fixed size of 41×41


$$SoH^\uparrow(G^*(R, W_i, st), G^*(R, W_{i+2}, st))$$

6.4 Results and Discussion

The results for the hot spot analysis are found in Fig. ?? and Fig. ?? for a comparison of G^* and the Focal G^* statistic. It can easily be seen that both versions produce similar results, but the focal versions produces a more differentiated picture for larger weight matrices. Small differences on a global scale are more pronounced on a regional scale and result in smaller and finer areas for hot spots. This enables the detection of additional hot spots and interesting areas which are most easily observable for the weight matrix of size 7×7 in the evening (Fig. ??). This enables the detection of significant deviations from the surrounding area. In contrast the Standard G^* statistic shows larger areas as important. Therefore, depending on the need of a planner, the Focal G^* statistic is more helpful to identify individual areas of interest whereas the standard G^* statistic gives a more broad overview. For the identification of IUHI this is quite important. As a city planer wishes to detect critical areas, it is important to detect not only general hot areas, but also those points where the most extreme differences in a local context exist. Finding those areas can help identify the underlying reasons or plan individual solutions.

Based on these images one can also see that the hot spots found by the Focal G^* statistic seem to be more stable. We compare their stability using only the SoH^\uparrow for lack of space. A plot of the results can be found in Fig. ?. Fig. ? compares the SoH^\uparrow between each increasing size of the weight matrix W. At a first glance, one can see that the typical implementation with a square weight matrix is the most unstable hot spot analysis, regardless of time of day. This is to be expected as the binary weights increase the dependence on the weight matrix. The use of a decreasing weight matrix perform better. As the more outlying data points get less weight, this reduces the dependence on the weight matrix and leads therefore to more stable results. Our proposed Focal G^* statistic achieves the most stable results in almost all cases. Only data points in a restricted region around the area of interest may influence the significance result. Through this restriction high values at key points gains more weight regardless of the

(a) Upward



(b) Downward




Figure 4: TODO

weight matrix and are therefore more independent of the weight matrix. This increases the stability. The decrease in stability for the largest weight matrices is most likely a result of the parametrization of the focal matrix. With increasing size of the weight matrix in relation to the focal matrix the value of each pixel is approaching to the mean of the area of the weight matrix. As can be easily seen from Def. 1 then the value for every pixel would be zero.

7 Conclusions and Future Work

In this work, we generalized the Getis-Ord statistic to deal with the problem of stability inherent in hot spot analysis. We identified possible underlying reasons for this instability: The weight matrix as well as the size of the study area. We developed a modified approach that deals with those two factors. The result is a modified G^* statistic called the *Focal G^* statistic*. It reduces the study area used for comparison into regions and achieves through this an increase in stability. To determine the effectiveness of our approach, we propose a stability metric for hot spots called *SoH*. To our knowledge no such metric existed before this work. The *SoH* computes the ratio of dependence of hot spots for different parametrizations of weight matrices. It enables to express the stability between each parametrization using single value restricted between zero and one. Based on this number one can decide which parametrization to use and researchers can compare the stability of their methods for unsupervised hot spot analyses. In particular, for temperature values one wishes to detect those areas which have high differences regardless of a particular parametrization. If a hot spot only appears for one parametrization, the information gained for general use is quite small and can even lead to an inefficient allocation of resources.

This research has several restrictions which have to be taken into account. First, we only tested the SoH^\uparrow metric. While we assume, based on our graphical analysis, that the SoH^\downarrow stability should be similar, we have no hard results. The results themselves are tested on two events in time for a fixed area of the city of Karlsruhe. We have not tested it on smaller or larger study areas, but we assume that the stability of the Focal G^* would stay the same whereas the stability of the G^* statistic would increase with a smaller study area and decrease with a larger study area. This follows the reasoning that the impact of a singular point increases with a decrease of the study area. To test this dependency is an interesting task for future work. The last restriction is the fixed size in this work of the focal matrix for the Focal G^* approach. We only tested one size in this work, but it is highly probable that the size of the focal matrix has an impact on the stability as could be seen in Fig. ???. While an overall trend can be seen in this work when the size of the weight matrix W and the focal matrix F are almost identical, the exact ratio is beyond the scope of this work. The optimal ratio as well as when the stabil-

ity suffers from a too similar size are interesting question for future work.

8 Acknowledgements

This work is part of the research project BigGIS (reference number: 01IS14012) funded by the Federal Ministry of Education and Research (BMBF) within the frame of the programme "Management and Analysis of Big Data" in "ICT 2020 – Research for Innovations". We thank the Nachbarschaftsverband Karlsruhe for the data of the thermal flight over Karlsruhe. R-packages used: raster [?], knitr [?]